QA Session

(SNLP tutorial)

Vilém Zouhar

21st July 2021

Perplexity

- How can we define it?
- Suppose that we generate a random sequence of digits S in the decimal system, where each digit in S is drawn from a uniform distribution. What is the perplexity of S in the following cases:
- - S is a sequence of length 5
- - S is a sequence of length 25

Word Sense Disambiguation

- Formal definition of the problem
- Naive Bayes for WSD
- Flip-Flop for clustering (???)

Conditional Random Fields

- First and second-order HMM
- Suppose you have a sequence of length M and tagset of size T, what would be the complexity of the normalization factor Z(x) in this case?
- Bayesian network, clicques

Naive Bayes

- Naive Bayes pseudocode
- Suppose you have a Naive Bayes classifier for topic categorization that is defined over a vocabulary of size V and a category label set of size C. How many parameters does this model have?
- Exercise 12 from 2020 exam

Naive Bayes - Code

for each class $c \in C$ # Calculate $P(c)$ terms

N_{doc} = number of documents in D

N_c = number of documents from D in class c

$logprior[c] \leftarrow \log \frac{N_c}{N_{doc}}$

$V \leftarrow$ vocabulary of D

$bigdoc[c] \leftarrow$ **append**(d) **for** d $\in$ D **with** class c

for each word w in V # Calculate $P(w|c)$ terms

$count(w,c) \leftarrow$ # of occurrences of w in $bigdoc[c]$

$$loglikelihood[w,c] \leftarrow \log \frac{count(w,c) + 1}{\sum_{w' \text{ in } V} (count(w',c) + 1)}$$

return *logprior*, *loglikelihood*, V

Naive Bayes - Code

Naïve Bayes

- Algorithm: Continuous-valued Features
  - Numberless values for a feature
  - Conditional probability often modeled with the normal distribution

$$\hat{P}(X_j \mid C = c_i) = \frac{1}{\sqrt{2\pi}\sigma_{ji}} \exp\left(-\frac{(X_j - \mu_{ji})^2}{2\sigma_{ji}^2}\right)$$

μ_{ji} : mean (avearage) of feature values X_j of examples for which $C = c_i$

σ_{ji} : standard deviation of feature values X_j of examples for which $C = c_i$

  - Learning Phase: for $\mathbf{X} = (X_1, \cdots, X_n)$, $C = c_1, \cdots, c_L$
    Output: $n \times L$ normal distributions and $P(C = c_i)$ $i = 1, \cdots, L$
  - Test Phase: Given an unknown instance $\mathbf{X}' = (a_1', \cdots, a_n')$
    - Instead of looking-up tables, calculate conditional probabilities with all the normal distributions achieved in the learning phrase
    - Apply the MAP rule to make a decision

Significance Testing

- How to use Chi-Square?

Compression

- Kraft's inequality and trees
- Optimal code length: $-\log_D p(w_i)$
- Encoding using a tree

Vector-Space Model

- Representation
- Retrieval (scoring vs bayes) + decision rule
- Classification

Jensen's Inequality

- For convex functions, opposite holds for concave ones
- Such as: x^2, log
- $WA(f(x_i)) \geq f(WA(x_i))$

Resources

- https://miro.medium.com/max/1400/1*neaBooRXSloZAz6A2XfHJA.png
- https://image.slidesharecdn.com/naive-bayes-150514165844-lva1-app6891/95/naive-bayes-15-638.jpg?cb=1431622795